



INTERNAL TECH DEMO & AGENTIC CASE STUDY

AI-Powered ATS Candidate Screener

Architecting an autonomous, evidence-driven AI screening agent with mathematical scoring rubrics, anti-hallucination guardrails, and automated PDF audit reports.



90% Faster

From 45 mins to 30 secs



Zero Skill Inflation

Strict deliverable verification



Audit-Ready Exports

Self-contained PDF & HTML twins

Presenter: Surendra Bharadwaj • **Skill Framework:** Claude/Antigravity Agentic Customization • **Target Domain:** SDET & QA Engineering

Why Traditional Resume Screening Breaks Down



Resume Keyword Padding

Candidates copy-paste modern keywords (Playwright, Agentic AI, MCP, RAG) into bare skills sections without having touched them in actual project deliverables.

- ▶ Standard ATS searches match raw strings blindly.
- ▶ No distinction between "hands-on expertise" and "passing buzzword".



Reviewer Fatigue & Inconsistency

Engineering managers and leads spend 30–45 minutes per profile reading lengthy 4–6 page PDFs, leading to subjective and inconsistent evaluation.

- ▶ Inconsistent rubric application across interviewers.
- ▶ Key qualification gaps missed before scheduling 1st technical round.



Expensive Interview Waste

Unqualified candidates pass initial HR screens, consuming high-value senior engineering hours only to fail on core coding or hard requirements.

- ▶ High cost per hire & bloated interview pipelines.
- ▶ Zero structured audit trail of why candidates were screened in/out.

How the Candidate Screener Agent is Built



Skill Folder Structure

Configured cleanly inside `.claude/skills/candidate-screener/` adhering to Antigravity / Claude Agentic Standards:

- ▶ `SKILL.md` : Master instruction set, operational rules, workflow triggers, and output schema.
- ▶ `references/job-description.md` : Active SDET role requirements, must-haves, good-to-haves, and experience gates.
- ▶ `references/rubric.md` : Mathematical 100-point rubric, status multipliers, and hard-override rules.
- ▶ `Multi-Format Ingestion` : Headless extraction across PDF, DOCX, DOC, TXT, and Markdown.



Agent Execution Pipeline

- STEP 1 Taxonomy Mapping:** Parse JD into weighted mandatory & good-to-have matrices.
- STEP 2 Document Ingestion:** Extract full text without hallucination or truncation.
- STEP 3 Experience Gate:** Check `Exp < 6 yrs` prior to scoring (Instant Fail if under).
- STEP 4 Evidence Audit:** Apply 4-tier status factor (Matched/Partial/Claimed/Missing).
- STEP 5 Delivery:** Render Markdown + HTML + Headless Edge PDF simultaneously.

The "No-Skill-Inflation" Evidence Engine

Unlike standard search tools, the agent demands **verifiable project deliverables**. A bare keyword in a skills tag cloud is NEVER awarded full points.

MATCHED (M) • 1.0x

Concrete Evidence

Skill is demonstrated in ≥1 project with tangible output. E.g.: *"Migrated test suite to Playwright JS, cutting runtime by 45%."*

PARTIAL MATCH (P) • 0.6x

Adjacent / Tangential

Adjacent tech used or passing project mention without full framework ownership. E.g., Web UI testing without explicit Cucumber BDD.

CLAIMED (C) • 0.3x

Resume Tag Padded

Appears in technical skills summary or tools bullet list, but **zero deliverables** exist in project history. Flags interview probe.

MISSING (U) • 0.0x

Zero Mention

Does not appear anywhere under the name, abbreviation, or verified industry synonyms.

🛡️ Anti-Hallucination Guardrail:

The agent is strictly constrained to the text extracted from the file. It cannot "infer" that a candidate knows Python just because they know Java, unless explicitly written and evidenced.

100–Point Scoring Model & Hard–Fail Overrides



Score Capacity Formula

Component	Max Weight	Description
Mandatory Skills	85 Pts	13 Core SDET skills (Languages, Playwright, Frameworks, UI/BDD, APIs, STLC, AI Testing)
Good-to-Have Bonus	10 Pts	Cloud (AWS/Azure), CI/CD, Pharma/Life Sciences domain, Low-code, Monitoring
Experience Fit	5 Pts	6–10 Yrs (5 pts), 11–12 Yrs (3 pts), >12 Yrs (2 pts), <6 Yrs (0 pts / Gate Fail)
TOTAL CAPACITY	100 Pts	≥80: Strong Fit 60–79: Potential Fit 40–59: Weak <40: Reject



Hard–Fail & Override Rules

Rule 1 — Experience Gate: Total Experience < 6.0 years → **Automatic Instant Rejection** prior to scoring.

Rule 2 — Core Language: Neither JS/TS nor Python evidenced → **Reject** regardless of score.

Rule 3 — AI Hard Gap: Both AI Testing & Agentic AI missing → **Cap score at < 60 and downgrade verdict tier.**

Rule 4 — Skills Padding: ≥ 3 claimed unevidenced skills triggers interview red flag and tier penalty.

Case Study: Screening 3 Real Candidate Resumes

Candidate	Exp	Score	Verdict	Primary Drivers & Override Rationale
Girish Joshi	10+ yrs (15.5)	64.5 / 100	PASSED POTENTIAL ⚠️	• Playwright + JS/TS evidenced at Luxoft • 8+ yrs Life Sciences/Pharma domain (Philips, Quintiles/IQVIA) • Evidenced AI Model Validation testing bullet
Tarun Killamsetty	9.5 yrs	59.0 / 100 (Capped from 67)	FAILED WEAK FIT (AI GAP) ❌	• Strong 9.5 yrs Playwright UI + Rest Assured API at Oracle CTMS • Triggered Rule #3 Override: Zero AI or Agentic experience • Capped score < 60 and downgraded from Potential to Weak
Shushil Kumar Koppu	5.3 yrs	0.0 / 100 (Disqualified)	FAILED GATE FAIL (<6Y) ❌	• Superb modern Playwright JS migration + Microservices API skillset • Triggered Rule #1 Gate: 5.3 yrs total experience vs 6.0 yrs hard cutoff

💡 Outcome:

Saved the team 2+ hours of interviewing. Selected **Girish Joshi** for Round 1 with targeted probe questions on API pipeline automation & Agentic architecture.

Tangible Impact & ROI for Engineering Teams



Massive Velocity Boost

- ▶ Screen 50+ profiles in seconds instead of spending full days.
- ▶ Instant PDF and Markdown reports generated autonomously.
- ▶ Fast-tracks top 5% candidates before competitors reach them.



Objective & Bias-Free

- ▶ Identical mathematical rubric applied to every profile.
- ▶ Immune to resume aesthetic embellishments or keyword stuffing.
- ▶ Clear, auditable justification for every score point awarded.



Targeted Interview Probes

- ▶ Automatically highlights "Claimed but Unevidenced" items.
- ▶ Gives interviewers pre-prepared questions for live technical rounds.
- ▶ Guarantees zero missed hard requirements before hiring decisions.

Blueprint: How You Can Build Custom Agents

You can replicate this exact agentic architecture for **any project or workflow** across QA, Development, DevOps, and Operations:

1

Define Instruction

`SKILL.md`

Define role, step-by-step procedure, triggers, and anti-hallucination rules.

2

Ground Truth Data

`references/`

Provide explicit standards, scoring rubrics, coding conventions, or checklists.

3

Tool Automation

`CLI / Python Scripts`

Hook parsers, linters, test runners, or headless Edge PDF printers into the flow.

4

Structured Outputs

`MD + HTML + PDF`

Produce clean tables, gap matrices, and exportable artifacts for stakeholders.

💡 Ideas for Your Next Team Agents:

- **Test Case Reviewer:** Audit BDD feature files against user stories.

- **PR / Code Quality Agent:** Enforce Playwright/TypeScript best practices.

- **Defect Triage Agent:** Classify bug logs against STLC severity rubrics.



WRAP-UP & DEMO

Ready to Build Next-Gen Agents?

The candidate screener proves how custom AI agents can eliminate tedious manual toil, ensure rigorous mathematical consistency, and elevate engineering standards.



Open Generated PDF Report



Read Full Implementation Guide

Press **Right Arrow** (→) or **Space** to navigate • **P** to print slides